



The Use of Artificial Intelligence Technology in Game Development

Haoxuan Wang*

University of Warwick

Coventry, United Kingdom

u5522898@live.warwick.ac.uk

Abstract

Artificial intelligence (AI) has become a transformative force in the gaming industry, reshaping development processes, player experiences, and operational management. This paper investigates both the potential and the challenges of applying AI technology in game development. It first examines how AI enhances game design through intelligent NPC behaviors, procedural content generation, and automated development tools, while also highlighting issues such as performance bottlenecks and repetitive content. Second, it analyzes AI's role in optimizing player experiences, focusing on affective computing, natural language processing, and dynamic difficulty adjustment, as well as the risks of misrecognition, over-personalization, and privacy concerns. Third, it evaluates AI's contributions to game operations, including anti-cheat systems, elastic server scaling, and matchmaking, alongside their inherent limitations. Based on these analyses, the paper proposes targeted solutions, ranging from robust adversarial reinforcement learning and content diversity frameworks to privacy-aware personalization and multi-level anti-cheat systems. Finally, it offers strategic recommendations for future development, including interdisciplinary collaboration, open platforms, and the integration of AIGC with procedural generation. The study provides both theoretical insights and practical guidance for advancing AI in the gaming industry.

CCS Concepts

• General and reference -> Cross-computing tools and techniques -> Reliability;

Keywords

Artificial Intelligence, Game Development, Procedural Content Generation (PCG), Dynamic Difficulty Adjustment (DDA)

ACM Reference Format:

Haoxuan Wang. 2025. The Use of Artificial Intelligence Technology in Game Development. In *2025 8th International Conference on Computer Information Science and Artificial Intelligence (CISAI 2025)*, September 12–14, 2025, Wuhan, China. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3773365.3773548>

*Corresponding Author



This work is licensed under a Creative Commons Attribution 4.0 International License. *CISAI 2025, Wuhan, China*

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1874-8/2025/09

<https://doi.org/10.1145/3773365.3773548>

1 Introduction

In recent years, with the improvement of living standards, the demand for entertainment has gradually increased, leading to the accelerated development of the gaming industry, with several new talents entering the field. Therefore, to stand out among these gaming companies, innovative technologies have become indispensable, and artificial intelligence (AI) is one key technology. The advanced applications of AI in other fields have also laid a solid foundation for its integration into game development. This paper will focus on the research problem of the current insufficient application of AI technology in the gaming industry and the need for its improvement. This research aims to explore how AI technology can enhance game development processes and player experiences, while also addressing the challenges that arise during its implementation. By focusing on the Resource-Based View and Institutional Theory, this study clarifies the strategic significance of AI technology as a key competitive resource in the industry, while also analyzing the challenges posed by the institutional environment during the implementation of technology and the corresponding strategies. It holds clear theoretical significance and practical guiding value. To comprehensively address these issues, the following sections analyze the applications, challenges, and solutions of AI in game design, player experience, and game operations, before proposing strategic directions for the industry's future development.

2 Analysis

This section will analysis the reasons for research questions from four aspects.

2.1 The Application of AI to Game Design

This section focuses on discussing AI's improvements to development processes and tools. In modern game development, artificial intelligence techniques simultaneously enhance developer productivity and player engagement through a range of key functionalities. First, NPC behavior modeling leverages deep reinforcement learning and behavioral feature analysis to endow non-player characters with adaptive and credible behaviors. It was demonstrated by OpenAI Five's strategic success in Dota 2 and rich behavior-characteristics frameworks [1]. Second, procedural content generation utilizes machine learning based PCG methods, including generative adversarial networks, to automatically generate diverse and repayable level designs [2]. Finally, AI-assisted scripting tools facilitate the automatic generation and optimization of game logic code, thereby reducing manual coding effort and accelerating iteration cycles.

2.2 Current Issues

Despite AI's powerful capabilities in optimizing player experiences, it still faces several major challenges. First, the current flexibility of NPC intelligent behaviors is insufficient. Deep learning-based NPC behavior models can generate unpredictable or anomalous actions when confronted with novel or edge-case scenarios. For example, an NPC may loop the same action indefinitely or exploit unintended game mechanics, thereby breaking immersion and causing players to feel annoyed or frustrated. Secondly, the real time computations required for high fidelity decision, making and dynamic level generation place considerable strain on the rendering pipeline. As shown in Figure 1, on low to mid-range hardware, this typically manifests as frame rate drops or extended loading pauses, interrupting gameplay flow and reducing user satisfaction [3]. The figure below plots the rendering time (in milliseconds per four frames) over a sequence of 140 game actions, clearly highlighting intermittent spikes where performance bottlenecks occur and frame rates drop. Third, procedural content generation systems, especially those trained on limited or homogeneous datasets, tend to reproduce similar layout patterns and art assets across levels. Without fine tuning diversity metrics or integrating human design oversight, the automatically generated environments can feel monotonous and repetitive, undermining players' motivation to continue exploring the game.

2.3 The Role of AI in Optimizing Player Experience

This section examines the role of AI from the perspective of player interaction and perception. AI can also amplify personalizing and deepening they're in-game experiences from the following aspects. Firstly, AI can leverage affective computing and natural language processing to detect and respond to players' emotional states and conversational cues, thereby deepening immersion. AI can also use a Dynamic Difficulty Adjustment (DDA) framework which will dynamically analyze player behavior and preferences in real time to adjust game difficulty, narrative branches, and reward pacing to sustain a continuous flow experience and high engagement [4].

However, the AI applications still present numerous challenges. Primarily, because current AI affective computing and natural language processing technologies are not yet fully mature, they cannot reliably distinguish sarcasm, humor, or culturally specific emotional expressions. As a result, NPCs may produce responses that are incongruent with the players' emotional state, thereby undermining player engagement. Another significant issue is over personalization, often referred to as the filter bubble effect. When AI systems continuously adjust game content based on player behavior, they can create an experience that is overly tailored to the individual, effectively limiting the range of possible interactions. This results in players being repeatedly exposed to the same types of challenges or narrative choices, restricting their ability to engage with and explore new game content. Moreover, AI-driven systems that collect and analyze real time data, such as players' speech, facial expressions, and behavioral cues, which can give rise to serious privacy concerns. While these technologies enhance immersion and personalization, they typically require continuous tracking of sensitive

player information, including emotional states and personal preferences, which may make players feel uncomfortable or perceive their privacy as being violated. This in turn poses additional ethical challenges for the gaming industry.

AI-driven operational tools significantly enhance backend efficiency and user experience through three core functionalities. First, in the realm of anti-cheat measures, the system integrates supervised learning models such as supporting vector machines and random forests to perform multimodal detection on player input sequences, network packets, and on-screen imagery, enabling the identification and interception of cheating behaviors within milliseconds, thereby preserving game fairness [5]. Second, to address peak online loads, AI can deploy predictive models that utilize historical traffic data and player login patterns, combined with reinforcement learning based scheduling policies or regression analysis, to proactively trigger elastic server scaling and ensure smooth operation under high concurrency conditions [6]. Finally, the intelligent matchmaking system, which is grounded in Bayesian skill rating and win rate prediction algorithms, computes players' skill levels and preferences in real time to dynamically form teams and allocate them into appropriate match pools, thereby not only enhancing competitive balance but also significantly boosting player satisfaction and retention, AI and player interaction and perception as shown in Figure 1 [7].

However, each core operational capability faces significant limitations. Currently, AI driven anti cheat systems trained on historical cheat samples may mistakenly classify expert players' high skill maneuvers as cheating. AI is also unable to detect certain novel cheating techniques, such as cheat programs that operate using DMA functionality. Moreover, elastic scaling models based on reinforcement learning or regression forecasting often exhibit oscillatory 'churn', which repeatedly scaling out and scaling in when confronted with sudden traffic spikes or sparse data, thereby wasting resources and potentially resulting in resource shortages at critical moments [8]. Finally, the intelligent matchmaking system suffers from a cold starts problem: new or inactive players lack sufficient historical data, resulting in initial matches that do not accurately reflect their skill levels. This mechanism is also vulnerable to experienced players deliberately smurfing with alternate accounts to disrupt fair matchmaking, thereby reducing long term retention.

Theoretical Operationalization: RBV and Institutional Drivers of AI Adoption. To move beyond merely naming the theories, we specify measurable constructs, concrete variables, and expected effects on AI adoption outcomes—defined along four dimensions: speed (time from prototype to live), scope (number of subsystems using AI: design, UX, ops), depth (share of gameplay minutes/events influenced by AI), and governance maturity (privacy-by-design, safety/fairness processes, audit readiness).

As shown in Table 1, RBV explains capability-driven heterogeneity: stronger data/MLOps and reusable AI assets enable faster, broader, and deeper deployment with better governance. Stakeholder Theory captures external expectations: community fairness/privacy demands, and platform requirements raise governance standards; with solid compliance readiness they also accelerate releases. Institutional Theory accounts for market/regulatory context: coercive regulation boosts governance but can slow teams lacking




Figure 1: AI and player interaction and perception

Table 1: Operationalizing RBV / Stakeholder / Institutional drivers of AI adoption

Lens	Construct	Observable variable (examples)
RBV	Data & MLOps capability	Event-type coverage, missing-data rate; model change lead time; rollback MTTR; feature-store coverage
RBV	AI talent & reusable components	AI engineers per 100 devs; avg years of experience; internal AI/PCG modules; reuse rate
Stakeholder Theory	Player/community legitimacy (fairness, privacy, explainability)	Reports per 1k matches; privacy complaints; community sentiment index
Stakeholder Theory	Platform/publisher/partner requirements	Platform review pass rate; SLA breach rate; third-party audits passed
Institutional Theory	Coercive pressure (regulatory intensity)	Applicable rules; DPIA coverage; fines/incident rate
Institutional Theory	Mimetic pressure (peer AI intensity)	Genre-level AI penetration; AI keywords in patch notes

readiness; mimetic pressure speeds diffusion yet risks “fast-but-shallow” adoption unless RBV capacities are strong.

3 Solutions

In response to the three major game design challenges identified above, this section will sequentially propose three targeted optimization strategies.

3.1 AI-Driven game design optimization solutions

To prevent NPC strategies from falling into repetitive loops or exploiting unintended game mechanics, a hybrid learning approach can be employed. Specifically, on top of the reinforcement learning framework, a set of base rules constraints is imposed. For example, if a particular sequence of behaviors is detected to repeat multiple

times within a short interval, the system triggers an escape mechanism to break the cycle. Concurrently, robust adversarial training methods pit the agent against a destabilizing adversary that applies perturbations to observations or environment dynamics, driving the policy to perform well under worst case conditions. In practice, NPCs trained using method called Robust Adversarial Reinforcement Learning (RARL) demonstrate stable and credible behaviors across diverse scenarios, effectively resisting exploitative strategies that would disrupt player immersion [9]. This combined strategy of applying formal design constraints alongside adversarial training both preserves policy creativity and adaptability and markedly improves NPC robustness, significantly reducing the need for post release debugging (as shown in Figure 2).




Figure 2: Hybrid learning for NPCs

To alleviate real time performance bottlenecks caused by high fidelity AI computations, non-critical tasks such as procedural content generation and complex decision-making models can be offloaded to background threads or dedicated edge inference servers and executed asynchronously to avoid blocking the main rendering loop. On the device side, techniques such as model pruning, quantization or knowledge distillation can be used to compress the model, reduce the number of parameters and lower inference latency without sacrificing accuracy. Additionally, a scheduler that partitions CNNs at runtime can be used when local CPU or GPU utilization is excessively high. The system will automatically offload heavy-weight network layers to edge or cloud inference; once the load decreases, it switches them back to the device, thus dynamically balancing performance and latency [10]. Meanwhile, combining shader caching, asynchronous resource streaming and techniques such as automatic tuning of batch sizes and thread affinity based on performance analysis further hides latency and optimizes hardware utilization. These strategies working in concert simplify AI workloads, maintain stable frame rates and deliver smooth gameplay on a variety of hardware platforms.

To address content homogenization in procedural content generation, designers can use a Quality Diversity (QD) framework, archive high performing generated instances, and apply diversity driven loss functions to optimize novelty and quality at the same time, avoiding the pursuit of a single objective, while balancing reconstruction fidelity and diversity goals, so that a single run can

produce a large number of distinct level segments [11]. Moreover, mixed-initiative PCG systems allow human designers to interactively guide generation: they can filter or adjust outputs in real time to prevent redundancy and align results with artistic intent. Finally, a hybrid creative system integrates human machine feedback loops, enabling designers to steer diverse metrics in real time, filter undesirable outputs, and adjust generation parameters, aligning procedural innovation with creative intent while preserving necessary unpredictability [12].

3.2 AI-Driven player experience optimization solutions

This section proposes the following three targeted solutions to the challenges. To reduce the high error rates in emotion and dialogue understanding, the model should first be trained on annotated corpora that includes sarcasm, humor, and culturally specific expressions, which can significantly improve the classifiers' ability to distinguish between literal and nonliteral meanings. At the same time, a cross-attention audio video fusion pipeline should be integrated so that it can jointly process vocal tone, lexical content, subtle facial expressions, and gesture dynamics to produce robust multimodal emotion embeddings, thereby improving the accuracy of affective analysis. In addition, by fine tuning context aware of Transformer architectures on multiturn dialogue logs, the system

can capture long range dependencies, discourse structure, and pragmatic nuances, thereby delivering more relevant NPC responses in complex conversational scenarios [13]. Finally, implement an optional player feedback system that allows players to label or correct NPC mis-responses, creating a continual learning mechanism so that model outputs increasingly align with player expectations and gradually eliminate misunderstandings. These measures will substantially enhance the realism of NPC interactions and deepen player immersion.

To prevent AI systems from trapping players in narrow information cocoons, the game engine can enforce controlled exploration quotas that replace a portion of personalized content with novel or underexplored material. Several studies have proposed strategies to increase recommendation diversity without degrading relevance. For example, diversified recommendation methods, which include multi criteria optimization and controlled randomness, can improve content coverage and serendipity by 30% to 50% [14], which means this prevents players from becoming trapped in narrow information cocoons, broadening their experiences and sustaining long term engagement without diminishing their core preferences. Additionally, the game can offer players an ‘exploration mode’ option, allowing them to choose whether they want recommendations of novel content and to set the proportion of such content, thereby achieving the ideal balance for each user. At the same time, analytics backends can apply differential privacy to aggregate player telemetry data such as engagement metrics and emotion scores into statistical summaries that protect individual anonymity [15]. For server-side AI processes that require greater computational power, encrypted inputs can be used without exposing raw data.

3.3 AI-Driven game operation optimization solutions

Building on the operational pain points identified earlier, this section outlines three countermeasures. To reduce both false positive and false negatives, modern anti-cheat systems should deploy multiple levels of defense. These measures work in concert to increase cheat detection coverage while maintaining low false positive rates, effectively addressing emerging threats. To provide new or inactive players with more accurate initial ratings, the system can arrange matches against AI bots in the first few games to assess their actual skill levels and set a reasonable range for skill differences in match-making. The two formulas below provide the calculation methods for the maximum allowable skill difference for an individual player and the total skill discrepancy of a team.

$$\begin{cases} |S_{p,i}^{\alpha} - S_{p',i'}^{\alpha}| \leq \tau \\ |S_{p,i}^{\beta} - S_{p',i'}^{\beta}| \leq \tau \end{cases} \quad (1)$$

For any two players p and p' in team α or team β , the absolute difference between their skill ratings estimated after the i^{th} and i'^{th} matches must not exceed a threshold τ [17].

$$\left| \sum_{p \in \alpha} S_{p,i}^{\alpha} - \sum_{p' \in \beta} S_{p',i'}^{\beta} \right| \leq \varphi \quad (2)$$

For two opposing teams α and β , let the total skill rating of each team be the sum of its members’ ratings, the absolute difference between these two totals must not exceed the threshold φ [17]. Based on the above formulas, to counter smurfing, a behavior detection pipeline should be implemented that aggregates features such as session duration, command sequence timing, win rate and item usage, and incorporates recent skill level anomaly scores to flag accounts whose gameplay patterns significantly deviate from their declared skill level.

Table 2 provides a concise overview of the main challenges, targeted solutions, and expected outcomes across game design, player experience, and operations. This consolidated view clarifies what to implement, how to implement it, and how to measure success for AI-enhanced game development.

4 Data Governance Framework and Future Development

To guide the industry’s long-term evolution, this section proposes several strategic development directions aimed at creating a game AI ecosystem that is more innovative, sustainable and socially responsible.

4.1 Data Governance Framework

This section specifies the controls required to collect, process, and operate player data for AI features in a way that is lawful, transparent, secure, and proportionate.

Table 3 condenses a GDPR-anchored governance flow into five merged phases with clear release gates. Collection & Consent: just-in-time notices, granular opt-ins, age-gating, per-field purpose tags (gates: 100% tagging, zero mis collection). Minimization & Ingestion: drop raw media, pseudonymize IDs, field allowlists, audited SDK inventory (gates: optional-field cap met, full SDK review). Training + PETs: DPIA/ROPA, differential privacy with published ϵ , on-device/federated by default (gates: leakage tests passed, DPIA coverage). Deployment/Monitoring: RBAC/ABAC, kill-switch, model cards, DP-protected dashboards, no raw PII in logs (gates: kill-switch present, alert SLAs). Rights/Retention: in-game privacy dashboard, field-level TTLs, cryptographic eraser (gates: DSAR ≤ 30 days, zero records beyond TTL).

4.2 Future Development

Progress in applying AI to games requires expertise that extends beyond computer science. By integrating insights from fields such as cognitive science, psychology, and human computer interaction, developers can create more realistic NPC behaviors and more intuitive systems. I recommend that game engine companies and universities establish joint research centers to accelerate the development of AI in gaming and advance the technology toward greater maturity. In addition, opening platforms that provide anonymized player data and standardized algorithm modules would promote collaborative innovation between industry and academia. Finally, incorporating social network analysis into game AI decision models can mitigate algorithmic bias and enhance the diversity of player interactions. To accelerate innovation and reduce duplicated effort, the industry should establish an open and auditable platform

Table 2: Summary of AI-driven solutions in game development

Area	Challenge	Proposed Solution	Expected Outcome / KPIs
Game Design	NPC loops / brittle behavior	Hybrid RL + rule constraints + RARL	More robust NPCs; fewer exploit bugs; improved immersion
	Real-time performance bottlenecks	Async offloading + model compression + dynamic partitioning	Stable FPS; reduced latency; wider hardware coverage
Player Experience	Emotion/dialogue misrecognition	Sarcasm-aware corpora + multimodal fusion + context Transformers	Improved NPC responses; deeper immersion
	Over-personalization (filter bubble)	Exploration quotas + diversified recommendations + exploration mode	Increased serendipity (30–50%); sustained engagement
Operations	Anti-cheat false ± / emerging hacks	Multi-layer defense + adversarial retraining + anomaly detection	Lower false positives; fairer play
	Elastic scaling churn	Predictive + reactive autoscaling with cost-aware penalties	Stable latency; fewer resource oscillations
	Matchmaking cold start & smurfing	Bot calibration + skill-diff thresholds + anomaly detection	Fairer initial matches; reduced smurfing; better retention

Table 3: GDPR-anchored controls for AI game data

Phase	Core controls (summary)	GDPR anchor	Release gate / KPI
Collection & Consent	JIT notices; explicit opt-in for sensitive; per-field purpose tags (default-off)	Arts. 5–6, 8, 13–14 [16]	100% fields tagged; consent logged; mis collection = 0
Minimisation & Ingestion	Drop raw media; pseudonymize IDs; vetted 3rd-party SDK list	Art. 5(1)(c), 28 [16]	Optional fields ≤ cap; 100% SDKs reviewed
Training Governance + PETs	DPIA/ROPA before high-risk use; differential privacy (publish ϵ); on-device/federated where feasible	Arts. 25, 30, 32, 35 [16] [17]	DPIA complete; ϵ recorded per release; leakage tests pass
Deployment, Access & Monitoring	RBAC/ABAC; kill-switch; model/system cards; no raw PII in logs; drift/bias alerts	Arts. 24, 25, 32 [16] [18]	Kill-switch present; zero-PII logs; alert SLA met
Rights, Retention & Deletion	In-game privacy dashboard (DSAR/opt-out/delete); field-level TTL; model retrain on deletes	Arts. 15–22, 5(1)(e) [16]	DSAR ≤ 30 days; records beyond TTL = 0

Table 4: Suggestions for Future Development of AI in Game Development

Area	Suggested Action	Expected Impact
Interdisciplinary Collaboration	Establish joint research centers with psychology, HCI, and cognitive science; promote industry–academia partnerships	More realistic NPCs; improved usability; faster maturity of AI in games
Open Platforms & Shared Resources	Create auditable platforms with anonymized datasets, standardized APIs/SDKs, testing tools, and resource leaderboards	Accelerated innovation; reduced duplication; higher transparency and trust
Integration of AIGC with PCG	Build pipelines combining AIGC for concept expansion with PCG for scalable generation; set clear generation rules	Scalable content creation; enhanced designer intent; reliable and safe production

that hosts anonymized player behavior datasets and provides standardized API and SDK modules for PCG, DDA, fair matchmaking systems, and anti-cheat capabilities. Moreover, each resource on the platform should provide compatibility with different systems, clear terms of use and privacy policies, appropriate testing tools,

and popularity leaderboards for resources of the same type, to ensure fair comparison and rapid iteration. This integrated platform will accelerate research and development, enhance comparability and trust, and consolidate fragmented datasets into reusable components deployable across multiple studios.

To achieve large scale production, studios should integrate AIGC (Artificial Intelligence Generated Content) with procedural content generation. A feasible pipeline is to first use AIGC to elaborate the creator's initial concept, then apply PCG to generate the corresponding content conditioned on the expanded specification. AIGC then performs self-checks for completeness and thematic relevance, after which the creator conducts a manual review, and the approved content is integrated into the game build and prepared for release. Therefore, a specific set of generation limits and game creation rules should be defined for AIGC, for example the number of levels. Table 4 summarizes the key strategic directions—collaboration, open platforms, and AIGC–PCG integration—that can guide the sustainable and innovative future of AI in game development.

5 Conclusion

Artificial intelligence has become the core driving force in game development, promoting sustainable industry growth by optimizing development workflows and enriching player experiences. At the same time, it is essential to address the cost, ethical, and talent challenges introduced by this technology. Looking ahead, open data and algorithm platforms, cross-disciplinary collaboration, and the integration of AIGC and PCG are essential to translating AI's potential into reliable, ethical, and scalable game experiences. From a theoretical standpoint, the findings reaffirm the Resource-Based View by positioning data assets, model pipelines, and MLOps capabilities as VRIN resources that are difficult to imitate and central to long-term advantage. Institutional Theory further explains heterogeneous adoption outcomes: studios are simultaneously shaped by regulatory pressures (privacy, online safety), normative expectations (fairness and anti-cheat), and mimetic forces (following market leaders), which together condition the speed and scope of AI deployment. In sum, realizing AI's promise in games requires not only stronger algorithms but also careful governance, privacy safeguards, and organizational capabilities. Studios that operationalize these elements in tandem are best positioned to deliver AI-enhanced experiences that are fun, fair, and future-proof.

References

- [1] Dharna, A. (2020). Control and Creation: Coevolution of Game Levels and Game-playing Agents (Master's thesis, Fordham University).
- [2] Liu, J., Snodgrass, S., Khalifa, A., Risi, S., Yannakakis, G. N., & Togelius, J. (2021). Deep learning for procedural content generation. *Neural Computing and Applications*, 33(1), 19–37.
- [3] Tufano, R., Scalabrino, S., Pascarella, L., Aghajani, E., Oliveto, R., & Bavota, G. (2022). Using reinforcement learning for load testing of video games. In *Proceedings of the 44th international conference on software engineering* (pp. 2303–2314).
- [4] Chen, S. (2024). Literature review of application of AI in improving gaming experience: NPC behavior and dynamic difficulty adjustment.
- [5] Sen, T. K. (2021). Interpersonal communication analysis with facial expression encoding and interactional modeling. University of Rochester.
- [6] Fregly, C., & Barth, A. (2021). Data Science on AWS. "O'Reilly Media, Inc".
- [7] Beal, R. J. (2022). Artificial intelligence in team sports (Doctoral dissertation, University of Southampton).
- [8] Adhikari, S. (2025). Real-Time Big Data Processing for Intelligent Transportation Systems: A Framework for Scalability. *Journal of Digital Transformation, Cyber Resilience, and Infrastructure Security*, 10(1), 1–10.
- [9] Lai, B. (2024). Learning by playing: a stock market simulation game with deep reinforcement Learning-powered NPCs.
- [10] Almeida, M., Laskaridis, S., Venieris, S. I., Leontiadis, I., & Lane, N. D. (2022). Dyno: Dynamic onloading of deep neural networks from cloud to device. *ACM Transactions on Embedded Computing Systems*, 21(6), 1–24.
- [11] Liu, Y., He, J., Liu, Y., Lin, X., Yu, F., Hu, J., ... & Dong, C. (2024). AdaptBIR: Adaptive Blind Image Restoration with latent diffusion prior for higher fidelity. *Pattern Recognition*, 155, 110659.
- [12] Javanmardi, E., Maresova, P., Xie, N., & Mierzwiak, R. (2024). Exploring business models for managing uncertainty in healthcare, medical devices, and biotechnology industries. *Heliyon*, 10(4).
- [13] Li, Y., Shen, X., Yao, X., Ding, X., Miao, Y., Krishnan, R., & Padman, R. (2025). Beyond single-turn: A survey on multi-turn interactions with large language models. *arXiv preprint arXiv:2504.04717*.
- [14] Wu, Q., Liu, Y., Miao, C., Zhao, Y., Guan, L., & Tang, H. (2019). Recent advances in diversified recommendation. *arXiv preprint arXiv:1905.06589*.
- [15] Kröger, J. L., Raschke, P., Campbell, J. P., & Ullrich, S. (2023). Surveilling the gamers: Privacy impacts of the video game industry. *Entertainment Computing*, 44, 100537.
- [16] European Parliament and Council. (2016). Regulation (EU) 2016/679 (General Data Protection Regulation). Official Journal of the European Union, L119, 1–88.
- [17] Dwork, C., & Roth, A. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4), 211–407.
- [18] National Institute of Standards and Technology (NIST). (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST Special Publication 1270.